

COMP90025 Parallel and Multicore Computing

Matrices and Gaussian Elimination

Aaron Harwood

School of Computing and Information Systems
The University of Melbourne

2019 Semester II

Matrix multiplication

An $n \times m$ matrix, A , has n rows and m columns,

$$A = \begin{bmatrix} a_{0,0} & a_{0,1} & \cdots & a_{0,m-1} \\ a_{1,0} & a_{1,1} & \cdots & a_{1,m-1} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n-1,0} & a_{n-1,1} & \cdots & a_{n-1,m-1} \end{bmatrix},$$


where $a_{i,j}$ is an element of A at the i -th row and j -th column.

An $m \times l$ matrix, A , can be multiplied by a $l \times n$ matrix, B , to produce an $m \times n$ matrix, C :

$$AB = C,$$

where

$$c_{i,j} = \sum_{k=0}^{l-1} a_{i,k} b_{k,j}.$$



For convenience, and without loss in generality, we can assume that all the matrices are square, $n \times n$. The sequential code explicitly shows the computations.

```
int A[n][n], B[n][n], C[n][n];  
:  
for(i=0;i<n;i++)  
    for(j=0;j<n;j++){  
        C[i][j]=0;  
        for(k=0;k<n;k++)  
            C[i][j]+=A[i][k]*B[k][j];  
    }
```

The number of computational steps required is $O(n^3)$.

Unlike prefix sum, there is no flow dependence on data and all computations are independent of one another, i.e. no multiplication $a_{i,k}b_{k,j}$ is used twice. Consequently a large number of processors can be used to cost effectively speed up the computation.

For $n \times n$ matrices and n processors, each row of the answer can be computed by a single processor. The total number of parallel computational steps is then $O(n^2)$.

For n^2 processors, each element of the answer can be computed by a single processor. The total number of parallel computational steps is then $O(n)$.


For n^3 processors it is possible to obtain a total number of parallel computational steps of $O(\log n)$. See that the summation is a reduction operation.

Usually the number of processors is much less than n when working with matrices of order $n \times n$. Although it may be possible to allocate n processes or threads, this approach leads to excessive overheads.

Consider a large matrix, A , of order $n \times n$, such that the matrix can be divided into s^2 submatrices, $A_{p,q}$, of order $m \times m$ where $m = \frac{n}{s}$. Thus $0 \leq p, q < s$.


For two large matrices, A and B , the matrix multiplication is equivalent to an element by element approach:

$$C_{p,q} = \sum_{k=0}^{s-1} A_{p,k} B_{k,q}.$$



Recursive subdivision

Consider matrices, A and B , of size $n \times n$ where $n = 2^t$. In this case the matrix can be conveniently subdivided until only single elements remain.



A

```

mat_mult(A, B, s){
    if(s==1) return C=A*B;
    s=s/2;
    P0=mat_mult(A00, B00, s);
    P1=mat_mult(A01, B10, s);
    P2=mat_mult(A00, B01, s);
    P3=mat_mult(A01, B11, s);
    P4=mat_mult(A10, B11, s);
    P5=mat_mult(A11, B10, s);
    P6=mat_mult(A10, B01, s);
    P7=mat_mult(A11, B11, s);
    C00 = P0 + P1;
    C01 = P2 + P3;
    C10 = P4 + P5;
    C11 = P6 + P7; }

```

Gaussian elimination

A system of equations in n unknowns is of the form:

$$a_{0,0}x_0 + a_{0,1}x_1 + \cdots + a_{0,n-1}x_{n-1} = b_0$$

$$a_{1,0}x_0 + a_{1,1}x_1 + \cdots + a_{1,n-1}x_{n-1} = b_1$$

$$\vdots$$

$$a_{n-1,0}x_0 + a_{n-1,1}x_1 + \cdots + a_{n-1,n-1}x_{n-1} = b_{n-1}$$

Written in matrix form this becomes:

$$Ax = b.$$

We are trying to find x that satisfies the equation.

Sparse matrices have many zeroes, while *dense* matrices have mostly non zero values.

Sparse matrices tend to admit more efficient methods for storing and computation, that may however also use approximations.

The *direct* method to solve a system of linear equations that have a dense coefficient matrix, A , is Gaussian elimination.

The objective of Gaussian elimination is to progressively “eliminate” variables from the system, until a single variable remains. This is the same method that you would use by hand.

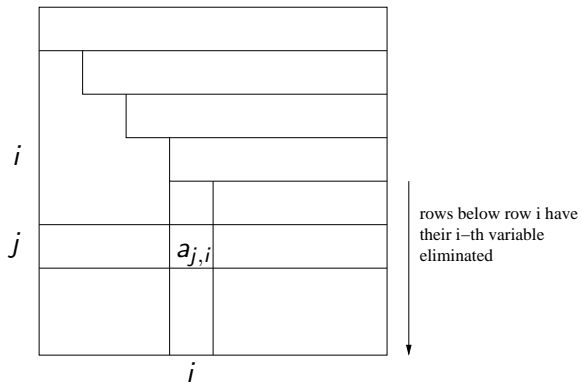
Given an equation in the system, say the i -th equation, we can eliminate an unknown from another equation, j , by multiplying equation i by a constant factor, $-a_{j,i}/a_{i,i}$, and adding the resulting equation to equation j . For example:

$$2x + 3y = 7 \text{ ;Equation 0}$$

$$5x + y = -4 \text{ ;Equation 1}$$

Multiplying Equation 0 by $-a_{1,0}/a_{0,0} = -5/2$ and adding to Equation 1 yields:

$$\frac{-13y}{2} = \frac{-43}{2} \Rightarrow y = \frac{43}{13}$$



One problem with Gaussian elimination is that when $a_{i,i}$ is zero or close to zero, the factor $-a_{j,i}/a_{i,i}$ is unstable.

In this case, the i -th row must be swapped, called *partial pivoting*, with the row, j , below i that has the greatest $a_{j,i}$. This does not affect the system of equations.

Partial pivoting can be done by finding the maximum in $\log_2 n$ steps. If all the variables are zero or close to zero then the variables in that column are already eliminated.

Ignoring the partial pivoting requirement, the sequential code is:

```
for(i=0;i<n-1;i++)
    for(j=i+1;j<n;j++){
        m=a[j][i]/a[i][i];
        for(k=i;k<n;k++)
            a[j][k]=a[j][k]-a[i][k]*m;
        b[j]=b[j]-b[i]*m;
    }
```

The time complexity is $O(n^3)$.

A parallel implementation can assign one row to each processor. With n rows, processor P_0 has row 0, processor P_1 has row 1 and so on.

The implementation has $n - 1$ phases. In the first phase P_0 broadcasts its row to all the other processors $\{P_i \mid i > 0\}$. These processors eliminate the first variable. In the second phase P_1 broadcasts its row (apart from the eliminated variable) to processors $\{P_i \mid i > 1\}$. These processors eliminate the second variable.

In the i -th phase, processor P_i broadcasts its remaining $n - i + 1$ variables (+1 is the b_i constant) to $n - i - 1$ remaining rows. Each processor must operate on $n - i + 1$ elements in each computational step.

Clearly some processors are idle for a significant fraction of the overall computation, e.g. P_0 becomes idle after the first phase. Partitioning the systems over a small number of processors is best done using *cyclic-striped partitioning*. This manages to keep all the processors busy for a greater period of time than block partitions.

| |
|-------------|
| processor 1 |
| processor 2 |
| processor 3 |
| processor 1 |
| processor 2 |
| processor 3 |
| processor 1 |
| processor 2 |